

# Lecture 8: Sorting in Linear Time

---

## Running time of sorting algorithms

Do you still remember what these statements mean?

- Sorting algorithm  $\mathcal{A}$  runs in  $O(n \log n)$  time.
- Sorting algorithm  $\mathcal{A}$  runs in  $\Omega(n \log n)$  time.

So far, all algorithms have running time  $\Omega(n \log n)$

- We didn't show this for heap sort, though it is true.

**Q:** Is it possible to design an algorithm that runs faster than  $\Omega(n \log n)$ ?

**A:** No, if the algorithm is comparison-based.

**Remark:** A comparison-based sorting algorithm is more general, e.g., the sorting algorithm implemented in C++ STL.

**Corollary:** Heap sort has running time  $\Theta(n \log n)$ .

## Lower bounds for problems

What does each of the following statements mean?

- Sorting can be done in  $O(n \log n)$  time.
  - There is a sorting algorithm that runs in  $O(n \log n)$  time.
- Sorting requires  $\Omega(n \log n)$  time.
  - Any sorting algorithm must run in  $\Omega(n \log n)$  time.

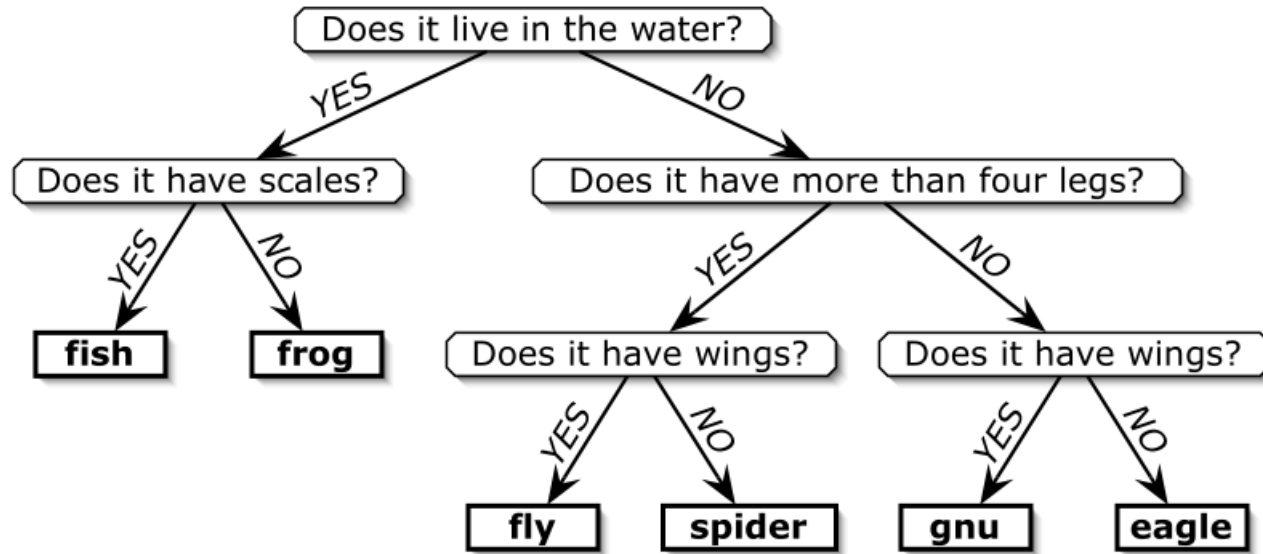
This is a very strong statement, which we unfortunately cannot prove.

- In fact, we have no lower bound higher than  $\Omega(n)$  for sorting.

But we can show an  $\Omega(n \log n)$  lower bound for comparison-based sorting algorithms

- More generally, in the decision-tree model of computation.

# The decision-tree model




A decision tree to choose one of six animals.

## An algorithm in the decision-tree model

- Solves the problem by asking questions with binary answers
- Cannot ask questions like "how many legs does it have?"
- The worst-case running time is the height of the decision tree.

## The decision-tree for binary search




**Theorem:** Any algorithm for finding a given element in a sorted array of size  $n$  must have running time  $\Omega(\log n)$  in the decision-tree model.

**Proof:**

- The algorithm must have at least  $n$  different outputs.
- The decision-tree has at least  $n$  leaves.
- Any binary tree with  $n$  leaves must have height  $\Omega(\log n)$ .

## The decision-tree for sorting

Sort  $\langle a_1, a_2, \dots, a_n \rangle$



**Theorem:** Any algorithm for sorting  $n$  elements must have running time  $\Omega(n \log n)$  in the decision-tree model.

**Pf:**

- The algorithm must have at least  $n!$  different outputs (there are so many different permutations).
- The decision-tree has at least  $n!$  leaves.
- Any binary tree with  $n$  leaves must have height  $\Omega(\log(n!)) = \Omega(n \log n)$ .

## Can we do better?

### Implication of the lower bound

- Have to use non-comparison based algorithms.
- Don't use worse-case analysis

### Integer sorting

- We will assume that the elements are integers from 0 to  $k$ .
- We will use both  $n$  and  $k$  to express the running time.
- Both  $n < k$  and  $n > k$  possible.

**Exercise.** Compare the following functions asymptotically:

$n, \log k, n + 1000, n + k, n \log k, \max(n, k), \min(n, k)$

# Counting sort

	1	2	3	4	5	6	7	8
<i>A</i>	2	5	3	0	2	3	0	3

	0	1	2	3	4	5
<i>C</i>	2	0	2	3	0	1

(a)

	0	1	2	3	4	5
<i>C</i>	2	2	4	7	7	8

(b)

	1	2	3	4	5	6	7	8
<i>B</i>							3	

	0	1	2	3	4	5
<i>C</i>	2	2	4	6	7	8

(c)

	1	2	3	4	5	6	7	8
<i>B</i>		0					3	

	0	1	2	3	4	5
<i>C</i>	1	2	4	6	7	8

(d)

	1	2	3	4	5	6	7	8
<i>B</i>		0				3	3	

	0	1	2	3	4	5
<i>C</i>	1	2	4	5	7	8

(e)

	1	2	3	4	5	6	7	8
<i>B</i>	0	0	2	2	3	3	3	5

(f)

*A*: input array

*B*: output array

*C*: counting array, then position array



## Counting sort

Counting-Sort( $A, B$ ):

let  $C[0..k]$  be a new array

for  $i \leftarrow 0$  to  $k$

$C[i] \leftarrow 0$

for  $j \leftarrow 1$  to  $n$

$C[A[j]] \leftarrow C[A[j]] + 1$

    //  $C[i]$  now contains the number of  $i$ 's

for  $i \leftarrow 1$  to  $k$

$C[i] \leftarrow C[i] + C[i - 1]$

    //  $C[i]$  now contains the number of elements  $\leq i$

for  $j \leftarrow n$  downto 1

$B[C[A[j]]] \leftarrow A[j]$


$C[A[j]] \leftarrow C[A[j]] - 1$

Running time:  $\Theta(n + k)$

Working space:  $\Theta(n + k)$

It is a stable sorting algorithm.

## Radix sort



Radix-Sort( $A, d$ ) :


for  $i \leftarrow 1$  to  $d$

    use counting sort to sort array  $A$  on digit  $i$

## Radix sort: Correctness

**Proof:** (induction on digit position)

- Assume that the numbers are sorted by their low-order  $i - 1$  digits
- Sort on digit  $i$ 
  - Two numbers that differ on digit  $i$  are correctly sorted by their low-order  $i$  digits
  - Two numbers equal on digit  $i$  are put in the same order as the input  $\Rightarrow$  correctly sorted by their low-order  $i$  digits



## Radix sort: Running time analysis

Q: How large should the “digits” be?

Analysis: Let each “digit” take values from 0 to  $b - 1$ .

- Counting sort takes  $\Theta(n + b)$  time.
- An integer  $\leq k$  has  $d = \log k / \log b$  such “digits”.
- So the running time of radix sort is  $\Theta\left(\frac{\log k}{\log b} \cdot (n + b)\right)$
- This is minimized when  $b = \Theta(n)$ .
- The running time is  $\Theta(n \log_n k)$ , which is  $O(n)$  when  $\log k = O(\log n)$

Q: When is this faster than  $\Theta(n \log n)$ ?

A: When  $\log_n k < \log n \Leftrightarrow \log k < \log^2 n$

Implementation:

- Choose  $n \leq b < 2n$  such that  $b$  is a power of 2.
- Use bit-wise operation for more efficient implementation.

## Summary of sorting algorithms

	Insertion sort	Merge sort	Quicksort	Heapsort	Radix sort
Running time	$\Theta(n^2)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n \log_n k)$
Working space	$\Theta(1)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(n)$
Randomized	No	No	Yes	No	No
Cache performance	Good	Good	Good	Bad	Bad
Parallelization	No	Excellent	Good	No	No
Stable	Yes	Yes	No	No	Yes
Comparison- based	Yes	Yes	Yes	Yes	No

# Bucket sort

## Average-case analysis

- Bucket sort is the only algorithm for which we do an average-case analysis in this course.
- We will assume that the elements are drawn uniformly and independently from a range.
  - We can assume the range is  $[0,1)$  without loss of generality.
  - They can be real numbers; while radix sort only works on integers

## Non-comparison based

- Bucket sort is non-comparison based
- It breaks the lower bound by breaking both conditions.


**Q:** Is it possible to design a comparison-based sorting algorithm with **expected or average running time** better than  $\Theta(n \log n)$ ?

**A:** No. (Proof omitted)

## Bucket sort

<i>A</i>
1 .78
2 .17
3 .39
4 .26
5 .72
6 .94
7 .21
8 .12
9 .23
10 .68

(a)



(b)

Worst-case:  $\Theta(n^2)$

**Bucket-Sort (*A*)**

let  $B[0..n - 1]$  be a new array

for  $i \leftarrow 0$  to  $n - 1$

    make  $B[i]$  an empty list

for  $i \leftarrow 1$  to  $n$

    insert  $A[i]$  into list  $B[\lfloor nA[i] \rfloor]$

for  $i \leftarrow 0$  to  $n - 1$

    sort list  $B[i]$  with insertion sort

concatenate the lists  $B[0], B[1], \dots, B[n - 1]$

together in order

## Average-case analysis of bucket sort

$$\begin{aligned} T(n) &= n + \sum_{i=0}^{n-1} n_i^2 \\ E[T(n)] &= E[n] + E\left[\sum_{i=0}^{n-1} n_i^2\right] \\ &= n + \sum_{i=0}^{n-1} E[n_i^2] \end{aligned}$$

**Claim:**  $E[n_i^2] = 2 - \frac{1}{n}$

**Pf:** Let  $X_{ij} = 1$  if  $A[j]$  falls into bucket  $i$ .

$$\begin{aligned} E[n_i^2] &= E\left[\left(\sum_{j=1}^n X_{ij}\right)^2\right] \\ &= E\left[\sum_{j=1}^n \sum_{k=1}^n X_{ij} X_{ik}\right] \\ &= E\left[\sum_{j=1}^n X_{ij}^2 + \sum_{1 \leq j \leq n} \sum_{\substack{1 \leq k \leq n \\ k \neq j}} X_{ij} X_{ik}\right] \\ &= \sum_{j=1}^n E[X_{ij}^2] + \sum_{1 \leq j \leq n} \sum_{\substack{1 \leq k \leq n \\ k \neq j}} E[X_{ij} X_{ik}] , \end{aligned}$$



## Average-case analysis of bucket sort

$$\begin{aligned} E[X_{ij}^2] &= 1^2 \cdot \frac{1}{n} + 0^2 \cdot \left(1 - \frac{1}{n}\right) \\ &= \frac{1}{n}. \end{aligned}$$

$$\begin{aligned} E[X_{ij} X_{ik}] &= E[X_{ij}] E[X_{ik}] \\ &= \frac{1}{n} \cdot \frac{1}{n} \\ &= \frac{1}{n^2}. \end{aligned}$$

$$\begin{aligned} E[n_i^2] &= \sum_{j=1}^n \frac{1}{n} + \sum_{1 \leq j \leq n} \sum_{\substack{1 \leq k \leq n \\ k \neq j}} \frac{1}{n^2} \\ &= n \cdot \frac{1}{n} + n(n-1) \cdot \frac{1}{n^2} \\ &= 1 + \frac{n-1}{n} \\ &= 2 - \frac{1}{n}, \end{aligned}$$

**Theorem:** When the inputs are drawn from a range uniformly and independently, bucket sort runs in expected  $\Theta(n)$  time.